

# Floorplanning Automation for Partial-Reconfigurable FPGAs via Feasible Placements Generation

Marco Rabozzi, *Student Member, IEEE*, Gianluca Carlo Durelli, *Student Member, IEEE*, Antonio Miele, *Member, IEEE*, John Lillis, *Member, IEEE*, and Marco Domenico Santambrogio, *Senior Member, IEEE*

**Abstract**—When dealing with partially reconfigurable designs on field-programmable gate array, floorplanning represents a critical step that highly impacts system's performance and reconfiguration overhead. However, current vendor design tools still require the floorplan to be manually defined by the designer. Within this paper, we provide a novel floorplanning automation framework, integrated in the Xilinx tool chain, which is based on an explicit enumeration of the possible placements of each region. Moreover, we propose a genetic algorithm (GA), enhanced with a local search strategy, to automate the floorplanning activity on the defined direct problem representation. The proposed approach has been experimentally evaluated with a synthetic benchmark suite and real case studies. We compared the designed solution against both the state-of-the-art algorithms and alternative engines based on the same direct problem representation. Experimental results demonstrated the effectiveness of the proposed direct problem representation and the superiority of the defined GA engine with respect to the other approaches in terms of exploration time and identified solution.

**Index Terms**—Field-programmable gate arrays (FPGAs), floorplanning, genetic algorithm (GA), local search, mixed-integer linear programming (MILP), partial reconfiguration (PR), simulated annealing (SA).

## I. INTRODUCTION

FIELD-PROGRAMMABLE gate array (FPGA) devices are nowadays widely employed in commercial and industrial appliances in many scenarios (e.g., telecommunication, automotive, high performance computing, and video and image processing), due to their reduced costs, good computational power, and high flexibility, since they can be reconfigured in order to change their functionality. Moreover, partial reconfiguration (PR) [1] has received a considerable attention in the recent years, since it even more enhances such flexibility by enabling the possibility to dynamically change only part of the modules at runtime while the rest of the system

Manuscript received November 16, 2015; revised March 8, 2016; accepted April 17, 2016. This work was supported by the European Commission through the H2020 FETHPC EXTRA Project under Grant 671653.

M. Rabozzi, G. C. Durelli, A. Miele, and M. D. Santambrogio are with the Dipartimento di Elettronica, Informazione e Bioingegneria, Politecnico di Milano, Milan 20133, Italy (e-mail: marco.rabozzi@polimi.it; gianlucacarlo.durelli@polimi.it; antonio.miele@polimi.it; marco.santambrogio@polimi.it).

J. Lillis is with The University of Illinois at Chicago, Chicago, IL 60607 USA (e-mail: lillis@uic.edu).

Color versions of one or more of the figures in this paper are available online at <http://ieeexplore.ieee.org>.

Digital Object Identifier 10.1109/TVLSI.2016.2562361

keeps working. Indeed, PR offers new opportunities such as the possibility to execute at different times more functionalities than the ones physically placeable on the device or the possibility to update or vary their implementations. In order to enable PR, two necessary conditions must hold: 1) the field programmable gate array (FPGA) device has to physically support the change of only a part of the configuration at runtime and 2) the companion design tools have to support the implementation of such reconfigurable systems. In this scenario, Xilinx [2] is the vendor presenting the most mature solution.

The role of floorplanning [3] in PR-based system design is even more prominent than to the standard FPGA design flow. In fact, while in the latter, this activity is mainly of interest for expert designers aiming at achieving advanced performance optimization; in the former, the implementation of a partially reconfigurable system forces to define the specific regions on the device fabric that will host the interchangeable functionalities. Therefore, floorplanning directly affects the feasibility and the performance of the final solution. However, it is a quite complex activity, since the area constraints for the reconfigurable regions have to meet specific placement requirements (reported in [1]) while covering a minimum amount of configurable resources that are needed by the modules reconfigured over time in each of the regions. Nevertheless, the internal architecture of FPGAs is becoming more and more advanced, exacerbating the floorplanning complexity. In fact, the homogeneous grid of configurable logic blocks (CLBs) is alternated, most of the time in an irregular way, with the columns of dedicated elements such as block RAMs (BRAMs) and digital signal processors (DSPs).

Commercial tool chains still support floorplanning through visual instruments, such as Xilinx Vivado [2] (which integrates the previous PlanAhead tool). However, the designer still has to manually define the shape and position of the reconfigurable regions, since the tool provides a limited automation on the region definition that generally leads to unfeasible solutions. Nevertheless, also other design flows for Xilinx FPGA devices proposed by the academia [4], [5] suffer from the same lack. On the other hand, several academic solutions have been presented in the literature to automate floorplanning [3], [6]–[12]. However, only few ones [6]–[8] take into account the requirements for PR, and, at the same time, accurately consider an

arbitrary distribution of heterogeneous resources within the device. Indeed, most of the algorithms consider only one of the two aspects, i.e., the PR requirements [9], [10] or the resource distribution [3], [11]–[13]. Finally, as it will be shown in this paper that such comprehensive approaches generate suboptimal solutions. Another relevant consideration that can be drawn on most of such automation solutions is the fact that they are actually unconnected from the real design flow; in fact, only few of these engines [4], [5], [14], [15] are tested on real circuits and synthesize their final outcome on a real board to check for feasibility.

In this paper, we present a novel floorplanning automation framework fully integrated with the Xilinx design flow. The framework exploits a direct representation of the problem based on the enumeration of the feasible placements that is able to abstract the computational complexity of floorplanning exploration while taking into account all the relevant constraints for PR on recent devices and metrics such as area consumption and aspect ratio. We show that, differently from the classical problem for VLSI design, enumerating a suitable subset of the feasible placements for each reconfigurable region is a viable approach and can also be efficiently automated by means of classical optimization algorithms. In a previous work [16], we proposed a preliminary version of the framework where the design space exploration was automated by means of a mixed-integer linear programming (MILP) formulation. Here, we propose a more complete and mature framework, featuring a new automation engine providing higher performance. In conclusion, we summarize our contributions as follows.

- 1) We accurately model all the constraints in the current PR guidelines [1].
- 2) We propose a direct formulation of the floorplanning problem based on a conflict graph of the feasible region placements described in terms of the actual coordinates on the fabric grid.
- 3) We propose a GA extended with a local search strategy exploiting the defined problem representation. The algorithm is able to speed up the identification of near-optimal solutions in a limited elaboration time.
- 4) Finally, we experimentally show the effectiveness of the proposed approach by comparing against various state-of-the-art solutions and alternative engines exploiting the same problem representation on both synthetic benchmarks and real case studies.

The remainder of this paper is organized as follows. Section II discusses the related work in the area, while Section III presents a formal description of the problem. Then, Section IV shows the proposed design flow, whose details are discussed in Sections V and VI, in which the feasible placement generation process and the floorplanning automation algorithm are presented, respectively. Finally, Section VII evaluates our approach on different problem instances, and Section VIII draws the conclusions. In addition, we report in the Appendix, a revisited presentation of the MILP formulation [16] that has been used as a baseline for the experimental evaluations.

## II. RELATED WORK

Several floorplanners for FPGAs have been proposed in the literature; however, most of them produce solutions that are either not compliant with PR requirements and guidelines [11]–[13], or only focus on a simplified device model, not capable of representing modern FPGAs lacking a uniform distribution of heterogeneous resources [9], [10].

One of the first algorithms that considers the heterogeneity of FPGA resources has been presented in [3]. The algorithm exploits simulated annealing (SA) over a slicing-tree representation and, subsequently, performs a compaction step to recover from unfeasible solutions and to improve the shapes of the modules. However, the resulting floorplan unlikely produces shapes that meet the PR requirements. Furthermore, the approach assumes the FPGA to have an homogeneous resource distribution, i.e., BRAM and DSP columns are homogeneously spaced within the device fabric. Based on this assumption, the algorithm divides the fabric grid in a set of homogeneous blocks having the same size and containing the same amount of resources for each resource type (DSP, BRAM, and CLB). However, this organization, which characterizes obsolete device families (such as Xilinx Virtex-II and Spartan 3), does not hold for the recent devices (e.g., Xilinx Virtex 6). Interesting aspects of such formulation are the irreducible realization list and the dominance relation, which have been successfully borrowed in our problem representation, as described in Section V.

A similar approach considering an heterogeneous FPGA device has been proposed in [13]; it consists of: 1) an SA algorithm exploring a sequence-pair representation of the solution and 2) a subsequent refinement of such a solution by means of a Min-Cost Max-Flow formulation, which alters the rectangular shapes of the reconfigurable regions. Due to this second phase, the approach in general does not satisfy PR requirements.

Another class of approaches [9], [10] introduces the time domain in the problem by handling the definition of reconfiguration operations together with the design of the floorplan. In [10], only logic blocks are taken into account while ignoring other types of resources available in the FPGA device. The work proposed in [9] considers both the partitioning of modules into reconfigurable regions and their floorplanning. During the partitioning phase, the algorithm assigns each of the modules to a reconfigurable region to minimize the wastage of resources over time. After partitioning, the resource requirements of the regions are known and the algorithm computes a floorplan by means of SA using a set of moves that preserve the PR constraints. Even though the approach considers heterogeneous resources, similar to [3], it assumes their regular and uniform distribution.

Differently from [9], other approaches [11], [12], called multilayer floorplanners, analyze together the various circuit configurations the system assumes in different instants of time. Their aim is to identify a floorplan, such that the common modules used in all the configurations are placed at the same position in all the circuit configurations. Such modules will represent the static area of the device, while the rest of the device is reconfigured as a whole.

Consequently, the reconfigurable part does not follow the Xilinx PR flow. Nevertheless, in [12], the device is assumed to have an homogeneous resource distribution as in [3].

A last class of floorplanners [6]–[8] considers both the PR constraints and an accurate description of the heterogeneous resource distribution. The work proposed in [6] stems from Parquet [17], the state-of-the-art fixed-outline floorplanner for VLSI design, and presents a nontrivial adaptation of the methodology to deal with partially reconfigurable FPGAs. The algorithm uses SA to perturb a floorplan representation that consists of a sequence pair augmented with a vector characterizing the aspect ratio of the modules. Moreover, to increase the probability to detect feasible floorplans, it implements smart moves to recover from solutions in which the resource requirements are not satisfied.

The approach devised in [7] characterizes the FPGA device in terms of minimal reconfigurable units [1] called tiles. Each tile spans multiple configurable frames on the horizontal direction and contains a specific type and number of resources. Thus, the resource requirements of the reconfigurable regions are translated in terms of tile requirements and a technique called Columnar Kernel Tessellation is applied to search for floorplans that minimize the overall estimated bitstream size. A postprocessing step moves the obtained areas on the vertical direction trying to locally improve the wire length without affecting the occupation of resources.

Even though [6] and [7] give better results than [9] in terms of wire length and area occupancy, respectively, [8] shows that the quality of their solutions can still be improved by means of analytic methods at the cost of a longer execution time. In particular, Rabozzi *et al.* [8] propose two algorithms both based on a compact MILP formulation. The first algorithm is meant to locally improve the quality of an initial feasible solution with a relatively small computational effort. Instead, the second algorithm is able, in principle, to explore the full solution space and to find provably optimal solutions. Unfortunately, both the algorithms require, to some extent, an initial feasible floorplan to achieve good final solutions. In our previous publication [16], we have demonstrated that the solutions achieved by the approaches proposed in [8] (and consequently in [6] and [7]) can be further optimized without additional time penalties.

A final aspect to be considered is the experimental validation of the proposed solutions. Actually, only in few approaches [4], [5], [14], [15], the produced results are synthesized on the target device to check their feasibility and the achieved performance in terms of maximum clock frequency. In particular, in [14], an in-depth analysis of the effects of module aspect ratio on the maximum achievable clock frequency is performed, while no automation strategies are presented. In [15], a similar analysis is performed by concluding that squared aspect ratios are preferable, and a very simple semiautomated floorplanner for pipeline designs based on a single chain of components is proposed. In [4], the floorplanning problem is tackled from a different perspective: first, the system is synthesized without any constraint, and then, an automated engine tries to identify a suitable set of placement constraints around the

TABLE I  
COMPARATIVE ANALYSIS OF PAST APPROACHES

Approach	FPGA Model*	Reconfig. aware	PR support	Experimental comparison**	Exp. verified
Cheng et al. [3]	Homo.			[17]	
Feng et al. [13]	Heter.				
Yuh et al. [10]	CLB	✓			
Montone et al. [9]	Homo.	✓	✓		
Singhal et al. [11]	CLB	✓			
Banerjee et al. [12]	Homo.	✓			
Bolchini et al. [6]	Heter.	✓	✓	[9], [13]	
Vipin et al. [7]	Heter.	✓	✓	[9]	
Rabozzi et al. [8]	Heter.	✓	✓	[6], [7]	
Lampricht et al. [14]	Heter.				✓
Neely et al. [15]	CLB	✓	✓		✓
Beckhoff et al. [4]	Heter.	✓	✓		✓
Yousuf et al. [5]	Heter.	✓	✓		✓
Rabozzi et al. [16]	Heter.	✓	✓	[6], [8]	
PA	Heter.	✓	✓	[16]	✓

(\*) Device models with a homogeneous resource distribution, heterogeneous one and considering only CLBs

(\*\*) The cell lists the approaches that have been tested and outperformed by the one of the current line

area used for placing and routing each module; unfortunately, the approach is tested with a single reconfigurable region and it is unlikely to work with a larger number of regions. Finally, in [5], an SA is directly integrated with the synthesis tool to implement each explored solution, even though such a strategy presents a huge cost in terms of elaboration time.

Table I recaps the characteristics of the existing approaches showing the supported features. It is worth noting that the most efficient approach has been proposed in [16]; in fact, it outperforms (possibly in an indirect way) most of the relevant previous solutions supporting PR [6]–[9], [13]. However, the weakness it presents is the lack of an experimental validation of the achieved solutions while not all the current PR constraints are taken into account. In this paper, we aim at proposing a novel floorplanning automation framework that, starting from the preliminary idea presented in [16], supports the peculiarities of modern FPGA devices and PR design flow, and features an even more efficient automation engine in terms of quality of the achieved solutions and elaboration time. Moreover, we also present an experimental validation of the approach by implementing real designs on an FPGA device.

### III. FLOORPLANNING PROBLEM DESCRIPTION

This section provides some relevant background on the floorplanning problem, in particular focusing on the Xilinx FPGA devices and the design rules of the related PR flow.

As shown in Fig. 1(a), the reconfigurable fabric of an FPGA device is organized in a set of columns of the resources of various types, that is,  $T = \{\text{CLB}, \text{BRAM}, \text{DSP}\}$ . The grid is also divided in quadrants, called clock regions, according to the structure of the clock tree and organization of the configuration memory. Based on the memory organization, the basic reconfiguration portion of the device grid, that we call tile, spans one clock region height and one resource width. Each tile contains a single type of resource depending on the position of the tile, and the amount of units depends on the type of resource. Thus, as in [7], we consider a more abstract model of the FPGA organization in terms of a grid of tiles. Finally, we also define a coordinate system on the grid of tiles,

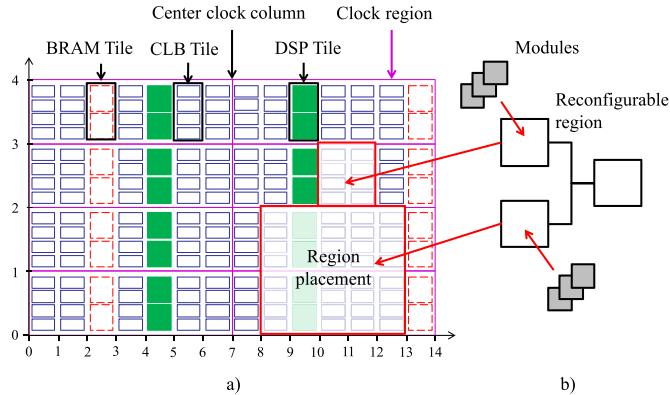


Fig. 1. Problem representation in terms of (a) reconfigurable FPGA device and (b) top level structural description of the system.

starting from the bottom-left corner. We denote with  $W$  and  $H$  the maximum values on the  $x$ - and  $y$ -axes, respectively.

According to the PR design guidelines, as shown in Fig. 1(b), the reconfigurable system is specified in terms of a structural description of interconnected  $N$  top components called reconfigurable regions.<sup>1</sup> Each region implements a partially reconfigurable unit in which it will be possible to load in a mutually exclusive fashion a set of modules implementing different functionalities. Thus, the reconfigurable region  $n$  presents resource requirements that depends on the hosted modules; for each resource type  $t$ , we denote the required amount as  $r_{n,t}$ . Moreover, the region is connected with the others and with the static part of the design (another component or set of components not featuring reconfiguration capabilities) by means of a set of interconnection buses, each one characterized by a width  $b$  in terms of number of wires. Do note that at the floorplanning stage, positioning of the connections of the wires among the region boundaries is not handled; therefore, a center-to-center interconnection model is here adopted, and the overall wire length is estimated using the classical half-perimeter wire length (HPWL) formula [12].

The goal of the floorplanning is to define a placement for each of the reconfigurable regions, in terms of rectangular shape and position on the FPGA resource grid. To this purpose, on the basis of the defined FPGA model, we denote with  $P$  the set of all possible placements that may be defined for the floorplanning of a single reconfigurable region

$$P = \{(x, y, w, h) \mid x, y, w, h \in \mathbb{N}, x + w \leq W, y + h \leq H\} \quad (1)$$

where  $x$  and  $y$  represent the coordinates of the bottom-left corner of the placement, while  $w$  and  $h$  define its width and height, respectively. Thus, the specific placement  $p$  can be characterized in terms of the available resource capacity, denoted as  $c_{p,t}$  (for each resource type  $t$ ), depending on the specific position and shape. It is worth noting that in some devices (e.g., the Zynq device), specific placements are forbidden, since they overlap with hard processors, static logic, or I/O blocks. We represent such placements with the subset  $S \subset P$ , which will be discarded during the floorplanning exploration. For a formal description of the floorplanning

<sup>1</sup>When clear from the context, we also refer to them simply as regions.

requirements, it is convenient to define a relation  $\perp$ , such that for  $p_1, p_2 \in P$ :  $p_1 \perp p_2$  if and only if the two placements overlap on at least a tile. The nonoverlapping relation  $\not\perp$  is simply defined as the complement of  $\perp$ :  $\not\perp = P \times P \setminus \perp$ .

To be feasible, a floorplan must assign a placement  $p_n$  for each region  $n$  and satisfy a set of PR requirements.

*REQ1:* Each assigned placement must contain at least the required resources for the corresponding region

$$\forall n \in N, t \in T : c_{p_n, t} \geq r_{n, t}. \quad (2)$$

*REQ2:* Each assigned placement must not be forbidden

$$\forall n \in N : p_n \notin S. \quad (3)$$

*REQ3:* The left and right boundaries of a placement  $p_n$  must be aligned to specific coordinates that prevent the splitting of interconnect resources [1] (VL and VR enumerate the valid left and right coordinates, respectively)

$$\forall p_n = (x, y, w, h) \mid n \in N : x \in VL \wedge x + w \in VR. \quad (4)$$

*REQ4:* CLB resources at both the sides of the center clock column must lie in the static part of the design, an assigned placement can cross the center column, but such resources are not available for the corresponding region

$$\forall t \in T : c_{(x_{clk}-1, 0, 2, H), t} = 0. \quad (5)$$

*REQ5:* Placements assigned to two different regions cannot overlap

$$\forall p_{n1}, p_{n2} \mid n1, n2 \in N \wedge n1 \neq n2 : p_{n1} \not\perp p_{n2}. \quad (6)$$

This list of constraints can be partitioned in two groups: REQ1–REQ4 are specifically related to the placement  $p_n$  for a single region  $n$ , while REQ5 rules the relative positions between different regions. Moreover, the first set can be summarized in a single definition by introducing a new set  $P_n$ , which represents all the feasible placements on the device for a reconfigurable region  $n$ .

In conclusion, the floorplanning problem can be stated as follows: given the sets  $P_n$  of feasible placements, a floorplan is a function  $f$  that assigns for each region  $n \in N$  a placement  $p \in P_n$ , such that there is no overlapping among the placements. More formally

$$f : n \in N \rightarrow p \in P_n \\ f(n_1) \not\perp f(n_2) \quad \forall n_1, n_2 \in N : n_1 \neq n_2. \quad (7)$$

#### IV. PROPOSED FLOORPLANNING FRAMEWORK

The structure of the proposed floorplanning automation framework and its integration in the Xilinx design flow are shown in Fig. 2. The design flow implemented in Xilinx Vivado consists in three main automated steps:

- 1) *synthesis*, which takes the input hardware description language (HDL) structural specification of the system and translates it in an intermediate netlist;

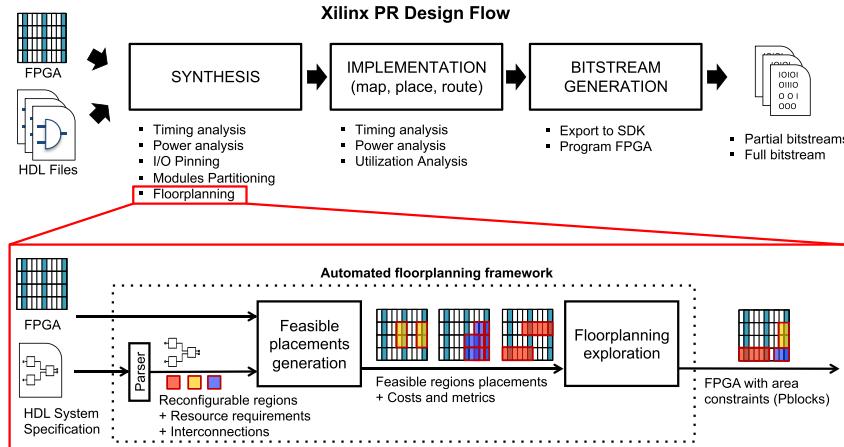


Fig. 2. Proposed floorplanning framework integrated within the Xilinx PR design flow.

- 2) *implementation*, which performs the place and route of the netlist on the selected FPGA device;
- 3) *bitstream generation*, which generates the final partial and complete configuration files.

Moreover, after each phase, a set of manual steps can be performed to define specific aspects and set parameters (such as the selection of the I/O pins or the floorplanning of the modules), while the obtained circuit can be analyzed by means of utility tools (for instance, to estimate the power consumption or the timing of the netlist).

In this scenario, the PR design flow is an enhancement of the standard flow able to handle the fact that several modules can be implemented in the same reconfigurable region. During the three phases, partial specifications, related to the modules and the top level of the design, are used for the synthesis and implementation of subcircuits and the subsequent generation of partial bitstreams. Within this scenario, the floorplanning is a manual activity executed before the implementation phase, immediately after the definition of the reconfigurable regions within the top-level specification.

The proposed floorplanning automation framework replaces the corresponding manual activity in the considered PR design flow. The framework takes in input the HDL structural specification of the system and translates it in an internal agile representation based on a graph. Moreover, it exploits synthesis reports to collect information on the resource requirements, which will be annotated on the system internal representation. The overall resource requirements of each reconfigurable region are computed as the maximum requirements among the various modules that will be hosted within the specific region. Finally, the framework takes in input the description of the considered FPGA device modeled as discussed in Section III. The output of the framework is the set of rules describing the floorplan solution, specified in the Xilinx constraint language to be imported in Vivado in order to continue with the subsequent implementation phase.

According to the formalization of the problem presented in the requirements REQ1–REQ5, the floorplanning automation framework is divided in two different phases.

- 1) Feasible placements generation that consists in building a conflict graph where nodes represent the union of

the  $P_n$  sets of possible feasible placements for each reconfigurable region  $n$  (REQ1–REQ4) and edges represent the overlapping among pairs of placements of different regions (REQ5).

- 2) Floorplanning exploration that selects a possible placement in  $P_n$  for each reconfigurable region  $n$ , such that there is no overlap among placements (REQ5) and an objective function specified by the designer is maximized.

We automated the two phases with different strategies according to their peculiarity and computational complexity. In particular, the first phase performs an exhaustive exploration for the definition of the conflict graph, since, as shown in Section V, the problem has a limited complexity. For the second phase, characterized by a considerably larger design space, the framework features an efficient exploration engine powered by a GA extended with a local search strategy. As shown in the experimental session, this strategy provides near-optimal solutions with a very limited execution time. Nevertheless, as shown in Section VII, the framework supports the integration of further automation strategies. The two phases are discussed in more detail in Section V and Section VI respectively.

## V. FEASIBLE PLACEMENT GENERATION

The first phase of the proposed framework is devoted to the definition of an abstract model called conflict graph that describes all the feasible placements for the various reconfigurable regions and the possible conflicts among pairs of placements. As shown in Fig. 3, the conflict graph contains a group of nodes for each reconfigurable region  $n$ , representing the overall enumeration of the feasible placements  $P_n$ , computed by fulfilling requirements REQ1–REQ4. Moreover, edges are used to represent conflicts between pairs of placements of two different regions, according to requirement REQ5.

It is worth noting that in the classical floorplanning for the VLSI design, it is commonly agreed that such a problem representation, based on the direct specification of all the possible region coordinates on the device grid, is extremely inefficient for automated optimization due to the huge solution space it defines. For this reason, past approaches have exploited various

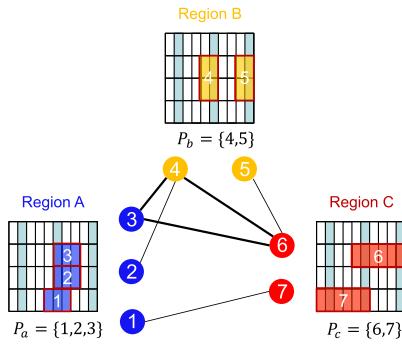


Fig. 3. Example of conflict graph.

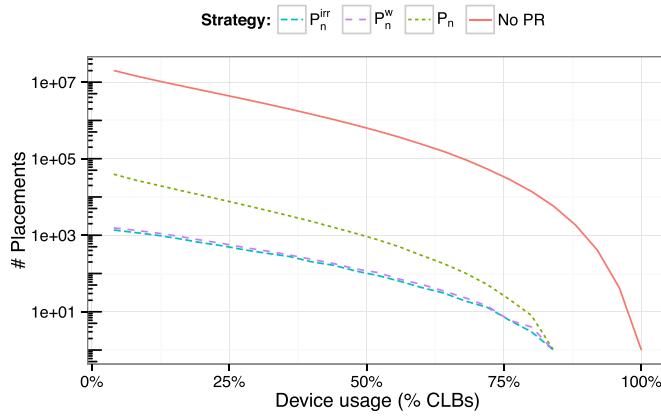


Fig. 4. Number of feasible placements for a single reconfigurable region.

indirect representations, such as slicing trees, sequence pairs, or other hierarchical tree-based representations [18]. At the opposite, the PR guidelines cause a considerable decrease in the number of feasible placements for a single reconfigurable region, thus allowing to effectively exploit such a direct representation of the placements. As an example, Fig. 4 shows the number of feasible placements generated for a single reconfigurable region when varying its resource requirement on a specific Xilinx XC7V585TFF6 device; the number of all feasible placements under PR constraints ( $P_n$ ) is two orders of magnitude smaller with respect to the number of all the possible placements that can be generated on the device without constraints (namely, *No PR* in Fig. 4). It is worth noting that only CLBs are considered as resource requirements, while taking into account that also other resource types would have even more decreased the number of placements.

The number of placements to be explored for the floor-planning can be even more reduced with respect to  $P_n$ , if we consider that in most of the cases, from an optimization point of view, it is not efficient to define placements larger than the minimal bounding boxes containing the required resources. In fact, as discussed in [3] and [6], by using the minimal bounding boxes, it is possible to reduce resource utilization, thus leaving space for additional functionalities and reducing the reconfiguration time. For this reason, we define a new set  $P_n^{\text{irr}}$  containing the irreducible placements of a region  $n$  as

$$P_n^{\text{irr}} = \{p \in P_n \mid \nexists p_2 \in P_n : p_2 \neq p \wedge p_2 \prec p\} \quad (8)$$

---

**Algorithm 1** Width-Reduced Placements Generation

---

```

1: for each  $n \in N$  do
2:    $P_n^w \leftarrow \emptyset$ 
3:   for each  $x \in VL$  do
4:     for  $y \leftarrow 0$  to  $H - 1$  do
5:       for  $h \leftarrow 1$  to  $H - y$  do
6:          $w \leftarrow \text{searchMinimalWidth}(x, y, h, n)$ 
7:          $v \leftarrow \text{validAspectRatio}(x, y, h, w)$ 
8:         if  $w > 0 \wedge v = \text{true}$  then
9:            $P_n^w \leftarrow P_n^w \cup (x, y, w, h)$ 

```

---

where  $\prec$  represents a containment relation between two different placements of the same region. More formally, given two placements  $p_1 = (x_1, y_1, w_1, h_1)$ ,  $p_2 = (x_2, y_2, w_2, h_2) \in P_n$ , we have  $p_1 \prec p_2$  if and only if  $x_1 \geq x_2$ ,  $y_1 \geq y_2$ ,  $x_1 + w_1 \leq x_2 + w_2$ , and  $y_1 + h_1 \leq y_2 + h_2$ . As shown in Fig. 4,  $P_n^{\text{irr}}$  allows to reduce the size of the conflict graph of about another order of magnitude compared with  $P_n$ .

Even if  $P_n^{\text{irr}}$  sets are well suited for optimizing resource occupation, we have experimentally noted that they may produce suboptimal results in terms of global wire length among various regions. For this reason, we have slightly relaxed the definition of  $P_n^{\text{irr}}$  to consider also the placements that are required to be minimal only with respect to the horizontal direction. For a formal definition of this set of placements, we consider a weaker containment relation  $\prec^w$ : given the two above placements  $p_1$  and  $p_2$ , we have  $p_1 \prec^w p_2$  if and only if  $x_1 = x_2$ ,  $y_1 = y_2$ ,  $h_1 = h_2$ , and  $w_1 \leq w_2$ . Thus, the corresponding width-reduced placements set for region  $n$  is defined as

$$P_n^w = \{p \in P_n \mid \nexists p_2 \in P_n : p_2 \neq p \wedge p_2 \prec^w p\}. \quad (9)$$

For the three defined sets, the following relation holds:

$$P_n^{\text{irr}} \subseteq P_n^w \subseteq P_n. \quad (10)$$

It is worth noting that the choice of reducing the placements only on the horizontal direction is suggested from the structure of current devices. Usually,  $H$  is much more coarse grained than  $W$ ; as an example, a Xilinx Virtex-5 XC5VLX110T is described using 8 rows and 62 columns of tiles ( $W = 62$  and  $H = 8$ ) [8]. On average, with respect to different CLB resource requirements, this relaxed strategy leads to 14% more placements with respect to  $P_n^{\text{irr}}$ , as shown in Fig. 4.

A last relevant issue related to the generation of the feasible placement is the aspect ratio, which is the ratio between the width and height of a placement. Indeed, as discussed in [14] and [15], extreme aspect ratios (e.g., lower than 1:5 or higher than 5:1) often lead to implementations with high routing congestion and low performance. This issue is mainly suffered on the vertical direction, since its axis is more coarse-grained. Thus, placements with such elongated shapes can be filtered during the placement generation process; then, among the available ones, higher cost can be attributed to the placements with extreme aspect ratios during the exploration phase.

Algorithm 1 automates the computation of the sets  $P_n^w$ . The procedure executes an extensive search of the possible placements by scanning all the valid coordinates of the device

starting from the bottom-left corner of the FPGA. For each point on the coordinate system, the algorithm considers all the possible placement heights that do not exceed the boundaries of the device and search for the minimal width needed to cover the required resources. Notice that depending on the resource requirements and on the presence of hard processors and static logic, the search for the minimal width can fail; in these situations, the corresponding placement is not generated. In the case of success, the *searchMinimalWidth* function returns the minimal width for the current starting point and height. If it is not possible to find a feasible placement for the given height and position, the method returns zero. At the same time, the *validAspectRatio* function is called to check if the current placement has to be discarded due to extreme aspect ratios.

From an asymptotic complexity point of view, *searchMinimalWidth* function is the most time-consuming operation in the innermost loop. By using a binary search and precomputing the resources occupied by each placement on the device, this function can be implemented with a  $\mathcal{O}(\log W)$  time complexity. Since the set VL, representing valid left coordinates for the regions, has a size proportional to  $W$ , *searchMinimalWidth* function is invoked  $\mathcal{O}(|N| \cdot H^2 \cdot W)$  times. In conclusion, Algorithm 1 has an overall time complexity of  $\mathcal{O}(|N| \cdot H^2 \cdot W \cdot \log W)$ . Notice that  $H$  and  $W$  are usually small numbers and the placement generation takes only a small amount of time compared with the overall optimization of the floorplan. Another observation, deriving directly from the algorithm, is that the number of width-reduced feasible placements for a region cannot exceed  $H^2 \cdot W$ . As discussed in Section VII, in real situations, the number of feasible placements is limited up to few thousands items per set, and therefore leads to manageable conflict graphs and to a placements generation time of few seconds.

## VI. FLOORPLANNING EXPLORATION

Once the feasible placements are generated, the second phase of the proposed floorplanning framework consists in the choice of the most suitable placement  $p$  for each region  $n$  among the available ones, such that: 1) all the selected placements do not overlap, meaning that no conflict edge exists among the pairs of such placements and 2) a specified objective function is optimized. The proposed floorplanning automation framework supports the integration of any automation engine capable of solving such an exploration problem by working on the defined conflict graph. It is worth noting that even if the conflict graph has a small size, the solution space, represented by the Cartesian product of the sets of feasible placements of each reconfigurable region, has a size that grows exponentially with respect to the number of regions, thus motivating the necessity of an efficient exploration engine.

In the preliminary formulation of the framework [16], we adopted an exploration engine based on an exact MILP model, whereas, within this paper, we designed and tested various heuristic methods in order to improve the floorplan exploration time. During our experimental sessions, we identified the GA engine extended with the steepest descent local search to be the most effective approach in finding

near-optimal solutions in a reduced amount of time. Therefore, we here present the GA engine in detail, while we refer the reader to the Appendix for a description of the MILP model that has been used as a baseline. From the experiments conducted with different placement sets, we also noticed that the width-reduced placements  $P_n^w$  offer the best tradeoff in terms of the size of the resulting solution space and quality of the achievable results. Hence, in the following discussion, we refer to sets  $P_n^w$ , even though the approach is still valid when other sets of placements, such as  $P_n$  or  $P_n^{irr}$ , are considered.

The proposed GA engine for automating the floorplanning exploration is based on the classical simple genetic algorithm formulation [19]. We defined a solution encoding and exploiting the enumeration of the feasible placements identified during the first phase of the floorplanning framework. More precisely, the chromosome is a linear vector, where each position represents a reconfigurable region  $n$ , and the contained value refers to the feasible placement  $p$  in the corresponding set  $P_n^w$ . Then, the standard crossover and mutation operators have been employed. The crossover operator cuts in a random point the chromosomes of two parent solutions and exchanges the second parts to generate two children, while the mutation operator replaces with a given probability the placement of a region with another randomly selected placement in  $P_n^w$ .

In order to evaluate the solution, we consider two different cost metrics:

- 1)  $A_{\text{cost}}$ , the cost directly related to placement selection;
- 2)  $W_{\text{cost}}$ , the cost deriving from the interregion wire length.

The first contribution can be easily computed summing the cost  $a_{p,n}$  associated with each placement  $p \in P_n^w$  that is selected for the current floorplan. As an example, the cost  $a_{p,n}$  can refer to the aspect ratio of the placement, amount of wasted resources, or wire length of a connection to a fixed I/O pin. On the other hand, the second metric estimates the interregion wire length using the HPWL formula. HPWL considers the wire connections concentrated in the center of the regions and measures the wire length using the Manhattan distance. In conclusion, the considered *fitness* function is a linear combination of the two defined metrics and an additional parameter  $\lambda$  able to handle unfeasible situations

$$\text{obj} = q_a \cdot \frac{A_{\text{cost}}}{A_{\max}} + q_{wl} \cdot \frac{WL_{\text{cost}}}{WL_{\max}} + \lambda. \quad (11)$$

In the formula,  $A_{\max}$  and  $WL_{\max}$  represent the maximum values that  $A_{\text{cost}}$  and  $WL_{\text{cost}}$  can assume, respectively; they are used to normalize the two contributions. Then,  $q_a$  and  $q_{wl}$  are user-defined weights. In particular, the fitness function first analyzes the feasibility of the solution only in terms of fulfillment of the nonoverlapping condition (REQ5), and then computes the cost according to the selected metrics. If the solution is unfeasible, a penalty value  $\lambda$ , defined as the number of pairs of regions that overlap, is summed to the objective value. In (11), we force  $q_a + q_{wl} = 1$ , so that valid floorplans are represented by  $0 < \text{obj} \leq 1$ , while  $\text{obj} > 1$  identifies unfeasible solutions. The goal of the  $\lambda$  factor is to enable the selection operator of the GA (we use the classical tournament selection) to rank unfeasible solutions in terms of the criticality of the constraint violation.

**Algorithm 2** GA Local Search

---

```

1: function IMPROVESOLUTION(solution)
2:   obj  $\leftarrow$  solution.evaluate()
3:   repeat
4:     oldObj  $\leftarrow$  obj
5:     for each n  $\in N$  do
6:       for each p  $\in P_n$  do
7:         obj'  $\leftarrow$  solution.evaluatePlacement(n, p)
8:         if obj'  $<$  obj then
9:           solution.setPlacement(n, p)
10:          obj  $\leftarrow$  obj'
11:   until obj  $<$  oldObj
12:   return solution

```

---

The choice of a simple solution encoding and operators has been driven by the possibility to directly manage the solution space, as motivated in Section V. However, as a drawback, we have noted during a preliminary experimental evaluation that such a GA engine is not able to obtain better performance than the preliminary MILP approach, since it generates too many unfeasible solutions. In fact, the direct problem formulation leads feasible solutions to evolve in unfeasible ones with a high probability. The main cause is the crossover operator that, due to its nature, applies global changes to each explored solution; at the opposite, the mutation operator that performs local moves has higher possibilities to make a feasible solution to evolve to another feasible one, that is, a neighbor in the solution space. For this reason, such an engine has been enhanced with a local search function, based on a steepest descent heuristic, which improves the current solution with iterated local modifications until no further improvement is possible. The strategy, shown in Algorithm 2, is invoked in the GA fitness function and guarantees to reach a local optimum from the input solution. We empirically demonstrated that the adoption of local search within GA leads to a hybrid approach able to converge faster toward global optima [20].

## VII. EXPERIMENTAL EVALUATION

The proposed floorplanning automation framework has been implemented in C++; GAlib [21] has been used for the GA engine. Within the following experimental sessions, we also integrated in the framework an SA engine, implemented in C++ by using the GNU Scientific library [22], and the preliminary MILP formulation [16]. For the other considered state-of-the-art approaches, the original algorithms provided by the related authors have been adopted, while all the MILP models have been solved using Gurobi 6.5. Finally, we integrated a graphical user interface (GUI) for analyzing and possibly modifying the floorplan solutions by means of a Web application in Javascript and HTML5.<sup>2</sup>

In the first experimental session, we performed an extensive testing campaign considering a large set of synthetic circuits aimed at demonstrating that the proposed GA engine outperforms the state-of-the-art approaches, whereas in Section II, we performed a more in-depth comparison of various engines (such as GA, SA, and MILP), exploiting the same direct problem representation proposed in this paper. Finally, we carried

<sup>2</sup>A preview of the GUI can be accessed at: <http://floorplacer.necst.it>.

out two real case studies to show that the framework generates feasible floorplanning solutions, possibly without any manual action of the designer, and moreover, it is able to improve the system performance. The three sessions are presented in Sections VII-A–VII-C. As a final note, all the experiments have been performed on a 2.2-GHz Intel Core Duo T6600 processor running a Linux operating system.

### A. Comparison With Respect to Past Approaches

The first experimental session considered the test suite of synthetic circuits from [8] targeted for the Virtex-5 XC5VLX110T device. This suite consists of 20 circuits with different area occupancy and number of reconfigurable regions; in particular, there are four circuits having a number of reconfigurable regions in the range {5, 10, 15, 20, 25}, while with respect to area utilization, there are five circuits for each fixed device occupancy in the range {70%, 75%, 80%, 85%}. It is worth noting that the maximum number of regions for the considered circuits has been set by taking into account that reasonably a reconfigurable system does not feature a larger number of reconfigurable regions. Moreover, we also made comparisons on some circuits from MCNC and GSRC suites, adapted as done in [6]; more precisely, we considered apte, xerox, hp, ami33, and ami49 targeted for a more recent Virtex-7 XC7K160T device.

In this first session, we compared our GA engine (dubbed PA-GA)<sup>3</sup> against the most efficient state-of-the-art approaches discussed in Section II, i.e., [6], the HO and O MILP-based algorithms presented in [8] and our preliminary MILP formulation [16] (dubbed as PA-MILP). Nevertheless, in order to perform the comparison, it was necessary to remove requirements REQ3 and REQ4 from the placement generation process, since previous approaches do not support them and their integration for [6] and [8] is not straightforward. Regarding the objective function, in this section, we only considered the overall wire length, since it has been noted to be the most challenging optimization goal. For each experiment, we executed ten runs of [6] and considered the best result as its final solution. According to the approach defined in [8], we run the HO approach by starting from some of the best solutions found in [6] (the ones within 10% from the best one) and, subsequently, O by using the final solution achieved by HO. For all the MILP formulations, the Gurobi solver execution time was limited to 1800 s, whereas for a PA-GA engine, we used a stopping criterion based on the elapsed time and the same time limit was applied. The elaboration was parallelized on all the available cores by using the *Threads* Gurobi setting and by running different processes for PA-GA with different random seeds. Notice, however, that the time limit does not take into account the time needed by PA-MILP and PA-GA for the generation of the feasible placements and the additional time to generate the initial solution for O [8]. Finally, PA-GA and PA-MILP used  $P_n^w$  as input.

Tables II and III show the results of this first experimental session, both in terms of execution time and quality of the

<sup>3</sup>The label PA, standing for proposed approach, here identifies all the engines based on the direct problem formulation proposed in this paper.

TABLE II  
RESULTS WITH DIFFERENT NUMBERS OF RECONFIGURABLE REGIONS

# RRs	Average wire length improvement w.r.t. [6]				Average execution time (sec)				
	HO[8]	O [8]	PA-MILP [16]	PA-GA	[6]	HO [8]	O [8]	PA-MILP [16]	PA-GA
5	6.99%	7.48%	7.46%	7.46%	11.0	13.3	84.1	32.5	1801.2
10	7.59%	11.65%	17.98%	19.07%	24.1	48.3	1848.3	1421.8	1801.4
15	8.88%	16.25%	34.03%	36.14%	41.1	74.6	1874.6	1802.1	1801.7
20	5.47%	14.80%	33.29%	39.17%	65.5	82.8	1882.9	1802.6	1801.9
25	5.67%	24.01%	40.72%	45.32%	94.0	119.1	1919.2	1803.0	1802.3

TABLE III  
RESULTS WITH DIFFERENT OVERALL DEVICE OCCUPANCY

Occupancy	Average wire length improvement w.r.t. [6]				Average execution time (sec)				
	HO [8]	O [8]	PA-MILP [16]	PA-GA	[6]	HO [8]	O [8]	PA-MILP [16]	PA-GA
70%	8.51%	17.40%	30.19%	32.15%	47.4	87.6	1559.0	1454.7	1801.7
75%	5.49%	18.89%	26.21%	28.74%	47.3	67.1	1514.2	1138.9	1801.7
80%	6.20%	12.86%	27.57%	30.73%	47.1	60.0	1508.1	1448.9	1801.7
85%	7.48%	10.20%	22.82%	26.10%	46.7	55.8	1505.9	1447.1	1801.7

TABLE IV  
APPROACHES COMPARISON ON DIFFERENT TEST CASES

Circuit	# RRs	Wire length					Execution time (sec)				
		[6]	HO [8]	O [8]	PA-MILP [16]	PA-GA	[6]	HO [8]	O [8]	PA-MILP [16]	PA-GA
apte	9	12789	12029	9682	5313	5206	21.75	147.93	1947.93	1801.87	1801.88
xerox	10	26589	25878	22974	12643	12813	25.53	307.97	2107.98	1802.36	1802.37
hp	11	12403	11796	11036	5568	5298	29.42	201.04	2001.05	1802.41	1802.45
ami33	33	172332	157583	130121	128538	104414	164.61	324.54	2124.57	1803.93	1803.98
ami49	49	55819	53178	36930	22669	18583	180.54	192.25	1992.26	1809.35	1809.47

achieved solutions. PA-GA was always able to find equivalent or better solutions than PA-MILP that in turn provided better results than [6] and [8]. Furthermore, the highest improvements were achieved for the most challenging circuits consisting of a high number of reconfigurable regions. In particular, for the test cases having 20 and 25 regions, PA-GA reduced the wire length of PA-MILP solutions by 8.2% on average while using the same amount of time. On the other hand, when considering the variation of resource usage (reported in Table III), as expected, the best results for both PA-GA and PA-MILP are obtained for the circuits with lower resource requirements even if there is no pronounced trend.

Table IV reports the results for the readapted MCNC and GSRC benchmarks. Results show that PA-MILP gives an improvement with respect to O that varies from 1.2% on ami33 to 49.5% on hp circuits, while with respect to the comparison between PA-MILP and PA-GA, the results are aligned to the synthetic benchmark trend. Indeed, for apte, xerox, and hp circuits, both PA-MILP and PA-GA provides a similar overall wire length, whereas when dealing with the bigger ami33 and ami49 circuits, PA-GA is able to improve PA-MILP solutions by 18.8% and 18.0%, respectively.

It is worth noting that for the synthetic circuits having five reconfigurable regions (Table II), PA-MILP is able to certify the optimality of the solutions and, thus, complete its execution before the given time limit. PA-GA, being a metaheuristic approach, cannot state if the identified solution is optimal,

and hence, the execution time of the algorithm depends on the time budget assigned. However, in these cases, we noted that PA-GA converge faster to the optimal solution than PA-MILP, while as the problem grows above the ten regions, no MILP formulation is able to reach the optimal solution in a reasonable time when the interregion wire length is considered. In fact, in our tests, we run the PA-MILP, i.e., the most efficient MILP formulation for several hours; however, after an initial very fast convergence to a near-optimal solution, the engine was not able to improve the solution or certify its optimality. This is related to the weak linear relaxation bounds provided by the MILP formulation with respect to the interregion wire length; the issue was only partially mitigated by using additional cuts to the model (see the Appendix). As an example, we report in Fig. 5 the graph representing the improvement of the best solution for the considered algorithms. We may note from Fig. 5 that PA-GA is the faster to evolve toward near-optimal solutions. This trend is representative for all the performed tests. Moreover, we noted that on average the PA-GA and PA-MILP tend to stabilize their solutions in <600 s, while after that, no relevant improvement is reported.

As a final note, the generation of the definition of the conflict graph model had a negligible impact on the overall execution time of the proposed algorithms. Indeed, in real situations, the size of the conflict graph is manageable; as an example, for the circuit having the highest number of

TABLE V  
PROPOSED APPROACHES COMPARISON WITH DIFFERENT NUMBER OF REGIONS

# RRs	Solutions improvements w.r.t. PA-MILP [16]								
	$q_a = 1.0, q_{wl} = 0.0$			$q_a = 0.5, q_{wl} = 0.5$			$q_a = 0.0, q_{wl} = 1.0$		
	PA-SA	PA-GAn	PA-GA	PA-SA	PA-GAn	PA-GA	PA-SA	PA-GAn	PA-GA
5	0.00%	-3.78%	0.00%	-0.77%	-11.95%	0.00%	-3.11%	-2.77%	0.00%
10	-4.31%	-23.02%	0.00%	-2.32%	-27.09%	0.18%	-8.59%	-29.16%	0.55%
15	-7.25%	-31.62%	-0.27%	-6.19%	-21.16%	5.28%	-12.69%	-78.73%	5.35%
20	-9.97%	-31.53%	-0.94%	-4.44%	-17.45%	7.09%	-15.29%	-30.36%	10.41%
25	-17.02%	-24.34%	-0.35%	-7.38%	-20.66%	9.42%	-3.75%	-8.78%	19.39%
30	-10.19%	-18.97%	-0.60%	-10.74%	-16.80%	8.95%	0.21%	-21.48%	21.86%

TABLE VI  
PROPOSED APPROACHES COMPARISON WITH VARYING RESOURCE USAGE

Usage	Solutions improvements w.r.t. PA-MILP [16]								
	$q_a = 1.0, q_{wl} = 0.0$			$q_a = 0.5, q_{wl} = 0.5$			$q_a = 0.0, q_{wl} = 1.0$		
	PA-SA	PA-GAn	PA-GA	PA-SA	PA-GAn	PA-GA	PA-SA	PA-GAn	PA-GA
70%	-9.63%	-17.24%	0.00%	-1.02%	-13.23%	7.50%	1.44%	-15.95%	10.67%
75%	-7.72%	-24.94%	-0.32%	-3.77%	-11.13%	5.98%	-3.16%	-21.95%	9.29%
80%	-6.37%	-19.66%	-0.28%	-11.66%	-28.70%	3.21%	-17.54%	-53.37%	7.95%
85%	-8.75%	-26.81%	-0.82%	-4.76%	-25.84%	3.94%	-9.52%	-25.12%	10.49%

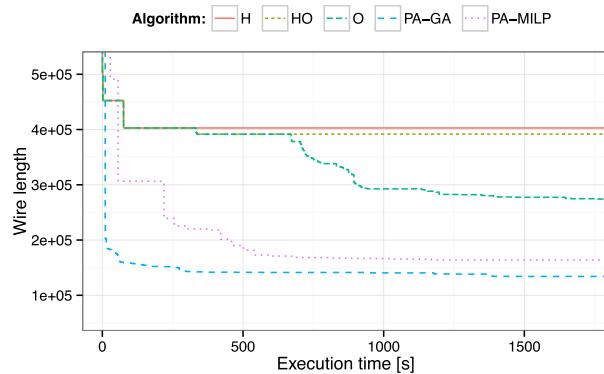


Fig. 5. Solution improvement over time for different approaches on the ami49 test case.

regions (ami49), the feasible placements generation process produced 146 446 nodes in less than 10 s.

### B. Analysis of Engines Based on the Proposed Representation

In a second session, we performed a more challenging comparison of PA-GA against other engines, exploiting the same direct problem representations. In particular, we considered: 1) PA-MILP; 2) an SA engine (called PA-SA), since it represents the classical approach for floorplanning strategies; and 3) a GA engine, exploiting the same solution encoding but without any local search strategy (dubbed as PA-GAn). The SA engine was defined on the basis of the standard SA algorithm and reusing the evaluation and mutation functions of the GA engine. In this experimental session, we considered PA-MILP as a baseline, since it has been the preliminary automating solution designed for the proposed framework. For this analysis, we included requirements REQ3 and REQ4, and we readapted the suite of the

first session to target a Virtex-7 XC7V585T device; we actually modified the resource requirements of the circuits according to the size of the new device, and we considered a new set of circuits with 30 regions. It is worth noting that the considered Virtex-7 device presents a higher heterogeneous distribution of resources and requires the introduction of more forbidden placements than the Virtex-5 one used in the first experimental session. Indeed, in some preliminary tests, we found that the considered state-of-the-art engines in [6] and [8] failed in finding any feasible solution on this device.

Moreover, the tests have been performed considering different settings of the objective function ranging from an optimization based only on the placements cost ( $q_a = 1.0$  and  $q_{wl} = 0.0$ ), one considering only the wire length ( $q_a = 0.0$  and  $q_{wl} = 1.0$ ), and finally a mixed objective function taking into account both metrics to the same extent ( $q_a = 0.5$  and  $q_{wl} = 0.5$ ). In order to perform a fair comparison of the approaches in terms of exploration efficiency, according to the discussion in Section VII-A, we fixed a limited time budget of 600 s that includes the time for the generation of the feasible placements. Except for the new time limit, the run settings for PA-MILP and PA-GA were as in Section VII-A, whereas we restarted PA-SA engine several times on the available cores using different random seeds until the available time budget elapsed.

Tables V and VI compare the obtained results according to the number of regions and device usage, respectively. For the problem instances consisting of five regions, the MILP approach and the GA were both able to find the optimal solution in all the cases, whereas the SA engine found optimal solutions only when the objective function was set to consider uniquely region placements cost. In general, the SA engine was almost never able to achieve better solutions than the MILP-based algorithm except for some problem instances consisting of large number of regions. This result

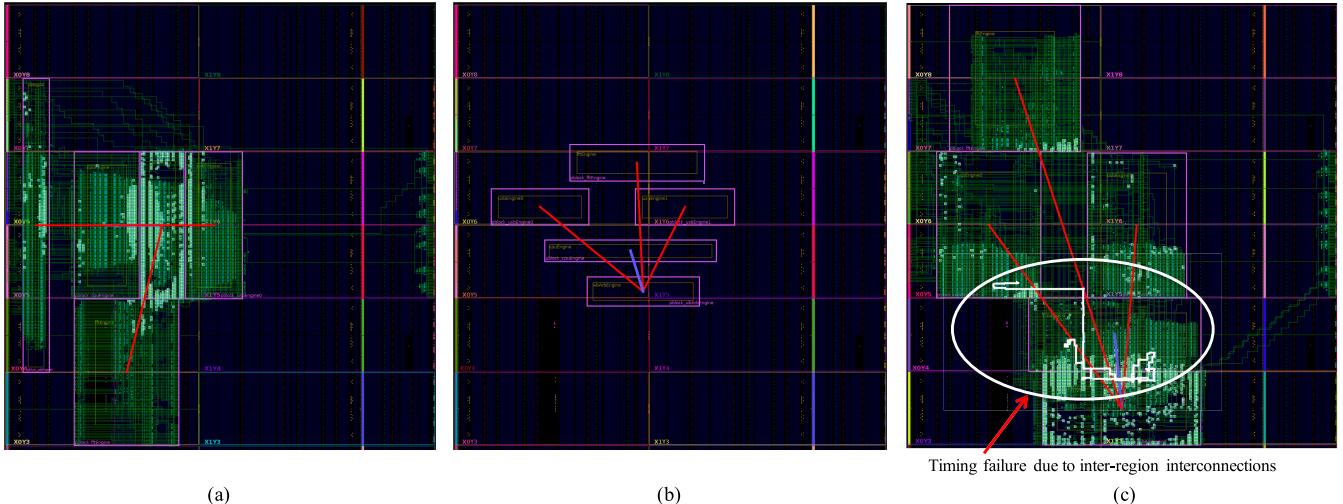


Fig. 6. Floorplans for the Xilinx case study. (a) Place and route of PA-GA floorplan at 100-MHz constraint. (b) Initial floorplan obtained with Vivado *place pblock*. (c) Place and route of manually adapted Vivado floorplan at 90-MHz constraint.

is quite interesting, since SA is the commonly used approach for automating the floorplanning exploration. On the contrary, we may conclude that it is not well suited for the defined problem representation due to the fact that many unfeasible solutions can be generated.

PA-GA proved to be an effective approach, leading to almost the same results for an optimization based on placement cost, while greatly outperforming the MILP engine when considering the most difficult problem, that is, the wire length optimization. This is especially highlighted when the circuits feature a large number of regions or a low resource usage. This improvement was mainly obtained exploiting local search within GA that allows to quickly explore a solution space consisting of local optimal solutions. Indeed, as it is possible to notice from Tables V and VI, PA-GAn provides results considerably far from its enhanced counterpart, since it spends a large amount of time exploring the unfeasible regions of the solution space, without the capability to recover from a feasible solution. It is worth noting that, in some experiments, PA-GAn was not able even to find an initial feasible solution; such situations were discarded from the results reported in Tables V and VI. Thus, we may conclude that the GA engine is the most promising solution for the proposed floorplanning framework.

### *C. Case Studies*

Finally, we validated the proposed approach on two real case studies to be implemented on a Virtex-7 XC7V585T device. The first case study is a readaptation of a Xilinx sample design (*project\_cpu\_virtex7*), consisting of five modules connected using a star topology, whereas the second case study is an in-house design of an image processing pipeline with reconfigurable components. As already noted in Section VII-B, the approaches proposed in [6] and [8] do not consider requirements REQ3 and REQ4, thus leading to potentially invalid floorplan solutions with respect to the subsequent place and route phase. Hence, we compared the results achieved by the proposed floorplanner with respect to the solutions manually

**TABLE VII**  
**RESOURCE REQUIREMENTS OF MODULES FROM THE XILINX CASE STUDY**

Module	LUTs	Registers	F7 Muxes	F8 Muxes	BRAMs	DSPs
cpuEngine	7440	3892	297	0	21	4
fftEngine	2837	1679	0	0	16	96
usbEngine0	6000	4699	259	81	36	0
usbEngine1	6080	4699	259	81	36	0
wbArbEngine	6800	1044	1959	172	0	0

designed by starting from the initial placements provided by the *place pblock* feature available in Vivado. This Vivado feature provides the user with suggestions on where to place the reconfigurable regions; however, the identified placements do not meet PR guidelines and require manual modifications.

Similarly to [15], the Xilinx case study has been readapted considering each of the five available modules as reconfigurable, thus leading to five distinct reconfigurable regions, each containing a single module. Notice, however, that this does not represent a limitation for the evaluation of the proposed approach, since the employed implementation flow is the same. The resource requirements of the reconfigurable regions are derived from the requirements of the corresponding modules (shown in Table VII) in which the number of LUTs was augmented by approximately 25% to ensure enough space for the insertion of proxy logic [1]. Furthermore, the number of interconnections among reconfigurable regions together with interconnections to I/O is summarized in Table VIII.

For the exploration, the objective function was set to consider wire length and resource consumption to the same extent in order to reduce both reconfiguration overhead and improve the possibility to meet timing constraints. The floorplan solution identified by PA-GA is shown in Fig. 6(a) together with the placed and routed circuit. Overall, the implementation phase was successful and the timing constraint requiring a 100 MHz frequency was met. On the other hand, Fig. 6(b) shows the initial solution provided by Vivado *place pblock*, and Fig. 6(c) the subsequent manually readapted floorplan. It is clearly visible that the designer has to perform

TABLE VIII

INTERCONNECTION MATRIX OF REGIONS FOR THE XILINX CASE STUDY

Region	fftEngine	usbEngine0	usbEngine1	wbArbEngine	I/O
cpuEngine	1	0	0	311	0
fftEngine	-	0	0	106	69
usbEngine0	-	-	0	118	69
usbEngine1	-	-	-	118	0
wbArbEngine	-	-	-	-	0

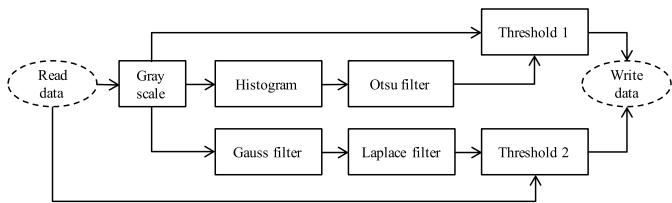


Fig. 7. Modules interconnections for the image processing case study.

a considerable and intrusive change of the solution proposed by Vivado; from our experience, we may report that such an activity requires around 2 h of time, while our automated engine requires a few minutes. Moreover, when considering the quality of the achieved solutions, this second one was not able to meet timing during implementation due to not optimized interregion interconnections. By lowering the timing constraint, it was possible to meet timing at 80 MHz; however, no place and route solution was found satisfying timing constraints with frequency equal to or higher than 85 MHz. Furthermore, the floorplan produced by the GA engine was able to reduce the overall size of the partial bitstreams by 25.7% with respect to the manual solution.

As a second case study, we realized a design in the context of image analysis consisting of seven different modules whose interconnections are shown in Fig. 7. In particular, the design is composed of two main computational pipelines that operate on a grayscaled image. The first pipeline includes the *histogram*, *Otsu filter*, and *threshold 1* modules, it binarizes the given image by applying the Otsu separation algorithm [23], while the second pipeline is configured to perform edge detection by exploiting the *Gauss filter*, *Laplace filter*, and *Threshold 2* modules. The modules were generated by using Vivado HLS, and AXI stream interfaces were used for communication. Furthermore, we implemented different alternative versions of each module within the design (employing different algorithms or having a different tradeoff between results accuracy and execution time) to exploit PR to switch from one to the other one; thus, seven reconfigurable regions were considered (one for each component within the original design). The resource requirements for each version of the various modules are listed in Table IX, whereas the requirements for the corresponding reconfigurable regions were obtained as in the previous case study.

Even in this scenario, the objective function was tuned to take into account the interregion wire length, regions aspect ratio, and wasted resources to the same extent. Within this case study, an analysis of the postimplementation results showed that the critical paths were represented by the internal

TABLE IX

RESOURCE REQUIREMENTS OF MODULES FOR THE IMAGE PROCESSING CASE STUDY

Region	Module	LUTs	FFs	BRAMs	DSPs
Laplace filter	LF_v1	628	332	64	2
Gauss filter	GF_3x3	807	465	64	0
	GF_3x3 float	881	809	32	5
	GF_5x5 float	815	760	32	5
Gray scale	GS_v1	334	238	64	4
Histogram	Hi_v1	256	180	1	0
	Hi_v2	104	87	1	0
Otsu filter	OF_v1	1205	1164	0	13
	OF_v2	726	517	0	2
Threshold 1	Th_v1	115	71	0	0
Threshold 2	Th_v1	115	71	0	0

TABLE X

MILP VARIABLES, SETS, AND PARAMETERS

Sets	
$N$	set of reconfigurable regions to floorplan
$P_n^w$	set of width-reduced feasible placements for region $n \in N$
$I$	set of interconnections between regions. Each element is a tuple of the form: $(n_1, n_2, b)$ where $n_1$ and $n_2$ are the regions involved in the interconnections and $b$ is its width
Parameters	
$W$	maximum value on the horizontal direction
$H$	maximum value on the vertical direction
$tileW$	the width of a tile within the FPGA
$tileH$	the height of a tile within the FPGA
$a_{n,p}$	cost associated to placement $p \in P_n^w$ for region $n \in N$
$q_a$	weight associated to the area cost
$q_{wl}$	weight associated to the wire length cost
$A_{max}$	maximum cost due to placements selection
$WL_{max}$	maximum cost related to global inter-region wire length
Variables	
$x_{n,p}$	binary variable set to 1 if and only if the placement $p \in P_n^w$ is selected for region $n \in N$
$cx_n$	$x$ coordinate of region $n \in N$ centroid
$cy_n$	$y$ coordinate of region $n \in N$ centroid
$dx_{n_1,n_2}$	horizontal distance between centroids of regions $n_1, n_2 \in N$
$dy_{n_1,n_2}$	vertical distance between centroids of regions $n_1, n_2 \in N$
$A_{cost}$	floorplan cost due to placements selection
$WL_{cost}$	floorplan cost related to global inter-region wire length

interconnections between the computational logic and the local BRAMs of the modules, whereas interregion interconnections were easily routed. Due to the peculiar characteristics of this design, both the PA-GA and the manually readapted floorplans were able to meet timing at 120 MHz and failing at a frequency equal to or higher than 125 MHz. Moreover, the floorplan produced by the GA engine was able to reduce the overall size of partial bitstreams of the manual solution from 9695 to 8815 kB, hence leading to a smaller reconfiguration time. Finally, similar to the previous case, the manual definition of the floorplan required about 2 h of activity.

TABLE XI  
MILP MODEL CONSTRAINTS AND OBJECTIVE FUNCTION

<b>Placements constraints</b>	
C1	$\sum_{p \in P_n^w} x_{n,p} = 1, \forall n \in N$
C2	$\sum_{n \in N, p \in P_n^w : p \perp (xp, yp, 1, 1)} x_{n,p} \leq 1, \forall xp \in [0, W-1], yp \in [0, H-1]$
<b>Wire length semantics</b>	
C3	$cx_n = \sum_{p=(xp, yp, wp, hp) \in P_n^w} x_{n,p} \cdot (xp + wp/2), \forall n \in N$
C4	$cyn = \sum_{p=(xp, yp, wp, hp) \in P_n^w} x_{n,p} \cdot (yp + hp/2), \forall n \in N$
C5	$dx_{n1,n2} \geq cx_{n1} - cx_{n2}, \forall n1, n2 \in N   n1 \neq n2$
C6	$dx_{n1,n2} \geq cx_{n2} - cx_{n1}, \forall n1, n2 \in N   n1 \neq n2$
C7	$dy_{n1,n2} \geq cyn_{n1} - cyn_{n2}, \forall n1, n2 \in N   n1 \neq n2$
C8	$dy_{n1,n2} \geq cyn_{n2} - cyn_{n1}, \forall n1, n2 \in N   n1 \neq n2$
<b>Cost functions definition</b>	
C9	$A_{cost} = \sum_{n \in N, p \in P_n^w} a_{n,p} \cdot x_{n,p}$
C10	$WL_{cost} = \sum_{(n1, n2, b) \in I} (dx_{n1,n2} \cdot tileW + dy_{n1,n2} \cdot tileH) \cdot b$
<b>Additional cuts</b>	
C11	$dx_{n1,n2} + dy_{n1,n2} \geq \sum_{p=(xp, yp, wp, hp) \in P_{n1}^w} x_{n1,p} \cdot \min\{wp/2, hp/2\} + \sum_{p=(xp, yp, wp, hp) \in P_{n2}^w} x_{n2,p} \cdot \min\{wp/2, hp/2\}, \forall n1, n2 \in N   n1 \neq n2$
<b>Objective function</b>	
OBJ	$\min \left\{ q_a \cdot \frac{A_{cost}}{A_{max}} + q_{wl} \cdot \frac{WL_{cost}}{WL_{max}} \right\}$

### VIII. CONCLUSION

This paper has proposed a novel floorplanning automation framework, compatible with the Xilinx tool chain and its PR flow. The framework considers a direct problem representation, which consists in an explicit enumeration of the possible placements of each region. The defined model allows to simplify the development of efficient floorplanning algorithms devoted to the optimization of different metrics such as aspect ratio, interregion wire length, and resource consumption. Various engines, based on an exact MILP formulation, GA, and SA heuristic approach, possibly enhanced with local search strategies, have been designed for automating the floorplanning activity. Such algorithms are experimentally evaluated with a challenging synthetic benchmark suite and real case studies. Experimental results demonstrated the effectiveness of the proposed direct problem representation and superiority of the defined GA engine with respect to the other defined strategies and the state-of-the-art approaches in terms of exploration time and identified solution.

### APPENDIX MILP FORMULATION

Within this appendix, we report the MILP model proposed in the preliminary version of the framework [16] and used within Section VII as a baseline for algorithms comparison. The variables, sets, and parameters of the formulation are listed in Table X, whereas the model constraints and objective function are summarized in Table XI.

In the proposed model, the binary variables  $x_{n,p}$  represent the current solution by stating which placement  $p$  is chosen for a given region  $n$ ; when considering the conflict graph, these variables state which specific node is selected for each region. Thus, a first class of constraints, dubbed as

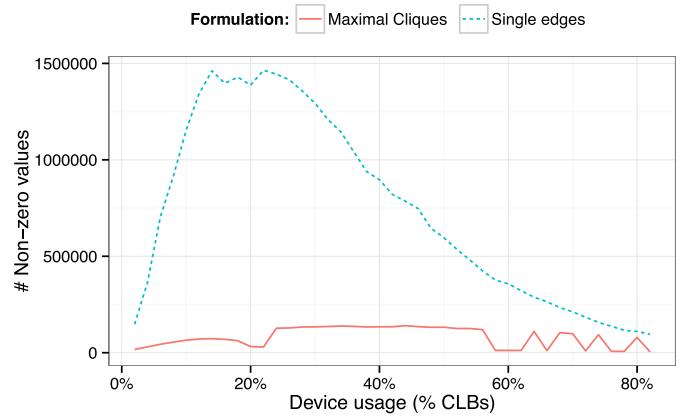


Fig. 8. Comparison of model size for the maximal clique and single edge MILP formulations on a problem consisting of two regions having equal resource requirements.

placements constraints, guarantees that a given solution is feasible: in particular, exactly one placement for each region must be selected (C1), and no pairs of placements connected by an edge can be selected, since they are overlapping (C2). It is worth noting that constraint C2 is defined on maximal cliques (i.e., fully connected subgraphs) in the conflict graph instead of single edges. This allows to reduce the size of the model while, at the same time, improves the linear relaxation bounds during the MILP solving process [24], so that the overall effect is a speedup in the MILP solver performance. Such a reduction in the problem size can be clearly seen in Fig. 8, by analyzing the number of edges between the placements generated for two regions when varying equally their requirements in terms of CLB; more precisely, Fig. 8 compares the number of nonzero terms within the constraint matrix of the complete MILP model using the single edges approach and the maximal cliques approach.

A second class of constraints, parameters, and variables are used to compute the cost OBJ of a given solution that is equivalent to the one used for the GA engine without the penalty contribution. Constraint C9 computes the  $A_{cost}$  metric by summing the cost  $a_{p,n}$  of each selected placement, while constraint C10 computes the global HPWL. The specification of C10 requires the introduction of variables  $cx_n$  and  $cyn$  to compute the coordinates of the centroid of each region  $n$  and variables  $dx_{n1,n2}$  and  $dy_{n1,n2}$  to compute the Manhattan distance between the centroids of each couple of regions  $n1$  and  $n2$ . Moreover, in order to guarantee the semantics of these variables, constraints C3–C8 are specified; constraints C3–C4 compute the coordinates of the centroids, while constraints C5–C8 ensure that the distances among regions cannot be less than expected.

An in-depth analysis of the formulation has shown that constraints C3–C8 give weak bounds when the linear relaxation of the MILP model is solved. For this reason, we introduced the additional cut C11, stating that the centroid distance of two regions has to be at least the sum of the distances to reach the centroids of the selected placements from their nearest borders. Indeed, a wire connecting the centroids of two placements has to cross at least one border for each of them.

## ACKNOWLEDGMENT

The authors would like to thank Xilinx, especially P. Lysaght, J. Wong, and R. Kong, for their support during this work.

## REFERENCES

- [1] Xilinx Inc. (2015). *Vivado Design Suite User Guide: Partial Reconfiguration*. [Online]. Available: <http://www.xilinx.com>
- [2] Xilinx Inc. *PlanAhead Design and Analysis Tool*, accessed on May 17, 2016. [Online]. Available: <http://www.xilinx.com/products/design-tools/planahead.html>
- [3] L. Cheng and M. D. F. Wong, "Floorplan design for multimillion gate FPGAs," *IEEE Trans. Comput.-Aided Design Integr. Circuits Syst.*, vol. 25, no. 12, pp. 2795–2805, Dec. 2006.
- [4] C. Beckhoff, D. Koch, and J. Torreson, "Automatic floorplanning and interface synthesis of island style reconfigurable systems with GoAhead," in *Proc. 26th Int. Conf. Archit. Comput. Syst. (ARCS)*, 2013, pp. 303–316.
- [5] S. Yousuf and A. Gordon-Ross, "DAPR: Design automation for partially reconfigurable FPGAs," in *Proc. Int. Conf. Eng. Reconfigurable Syst. Algorithms (ERSA)*, 2010, pp. 1–7.
- [6] C. Bolchini, A. Miele, and C. Sandionigi, "Automated resource-aware floorplanning of reconfigurable areas in partially-reconfigurable FPGA systems," in *Proc. Int. Conf. Field Program. Logic Appl. (FPL)*, Sep. 2011, pp. 532–538.
- [7] K. Vipin and S. A. Fahmy, "Architecture-aware reconfiguration-centric floorplanning for partial reconfiguration," in *Proc. Int. Conf. Reconfigurable Comput., Architect., Tools Appl. (ARC)*, 2012, pp. 13–25.
- [8] M. Rabozzi, J. Lillis, and M. D. Santambrogio, "Floorplanning for partially-reconfigurable FPGA systems via mixed-integer linear programming," in *Proc. Int. Symp. Field-Program. Custom Comput. Mach. (FCCM)*, May 2014, pp. 186–193.
- [9] A. Montone, M. D. Santambrogio, D. Sciuto, and S. O. Memik, "Placement and floorplanning in dynamically reconfigurable FPGAs," *ACM Trans. Reconfigurable Technol. Syst.*, vol. 3, no. 4, Nov. 2010, Art. no. 24.
- [10] P.-H. Yuh, C.-L. Yang, and Y.-W. Chang, "Temporal floorplanning using the three-dimensional transitive closure subGraph," *ACM Trans. Design Automat. Electron. Syst.*, vol. 12, no. 4, Sep. 2007, Art. no. 37.
- [11] L. Singhal and E. Bozorgzadeh, "Multi-layer floorplanning on a sequence of reconfigurable designs," in *Proc. Int. Conf. Field Program. Logic Appl. (FPL)*, Aug. 2006, pp. 1–8.
- [12] P. Banerjee, M. Sangtani, and S. Sur-Kolay, "Floorplanning for partially reconfigurable FPGAs," *IEEE Trans. Comput.-Aided Design Integr. Circuits Syst.*, vol. 30, no. 1, pp. 8–17, Jan. 2011.
- [13] Y. Feng and D. P. Mehta, "Heterogeneous floorplanning for FPGAs," in *Proc. Int. Conf. VLSI Design*, Jan. 2006, pp. 257–262.
- [14] J. Lamprecht and B. Hutchings, "Profiling FPGA floor-planning effects on timing closure," in *Proc. Int. Conf. Field Program. Logic Appl. (FPL)*, Aug. 2012, pp. 151–156.
- [15] C. E. Neely, G. Brebner, and W. Shang, "ReShape: Towards a high-level approach to design and operation of modular reconfigurable systems," *ACM Trans. Reconfigurable Technol. Syst.*, vol. 6, no. 1, pp. 5:1–5:23, May 2013.
- [16] M. Rabozzi, A. Miele, and M. D. Santambrogio, "Floorplanning for partially-reconfigurable FPGAs via feasible placements detection," in *Proc. Int. Symp. Field-Program. Custom Comput. Mach. (FCCM)*, May 2015, pp. 252–255.
- [17] S. N. Adya and I. L. Markov, "Fixed-outline floorplanning: Enabling hierarchical design," *IEEE Trans. Very Large Scale Integr. (VLSI) Syst.*, vol. 11, no. 6, pp. 1120–1135, Dec. 2003.
- [18] S. H. Gerez, *Algorithms for VLSI Design Automation*, vol. 8. New York, NY, USA: Wiley, 1999.
- [19] D. E. Goldberg, *Genetic Algorithms in Search, Optimization, and Machine Learning*. Reading, MA, USA: Addison-Wesley, 1989.
- [20] W. Wan and J. B. Birch, "An improved hybrid genetic algorithm with a new local search procedure," *J. Appl. Math.*, vol. 2013, Aug. 2013, Art. no. 103591.
- [21] M. Wall. (2007). *GAlib A C++ Library of Genetic Algorithm Components*, accessed on May 17, 2016. [Online]. Available: <http://lancet.mit.edu/ga/>
- [22] (2014). *GNU Scientific Library*. [Online]. Available: <http://www.gnu.org/software/gsl>
- [23] N. Otsu, "A threshold selection method from gray-level histograms," *Automatica*, vol. 11, nos. 23–27, pp. 285–296, 1975.
- [24] S. Rebennack, "Stable set problem: Branch & cut algorithms," in *Encyclopedia of Optimization*, C. A. Floudas and P. M. Pardalos, Eds. New York, NY, USA: Springer, 2009, pp. 3676–3688.



field-programmable gate

**Marco Rabozzi** (S'14) received the M.Sc. degree in computer science from The University of Illinois at Chicago, Chicago, IL, USA, in 2014, and the master's degree in computer science and engineering from the Politecnico di Milano, Milan, Italy, where he is currently pursuing the Ph.D. degree.

His current research interests include combinatorial optimization and computer-aided design, with a focus on mixed-integer linear programming, algorithms for resource constrained scheduling problems, and automated floorplanning tools for array devices.



**Gianluca Carlo Durelli** (S'13) received the B.Sc. and M.Sc. degrees in computer engineering, and the Ph.D. degree in computer science from the Politecnico di Milano, Milan, Italy, in 2009 and 2012, 2016, respectively.

He is currently a Research Affiliate with the Politecnico di Milano. His current research interests include embedded systems, computer architectures, and operating systems, with a particular focus on runtime resource management.



**Antonio Miele** (M'12) received the master's degree in computer engineering from the Politecnico di Milano, Milan, Italy, the M.Sc. degree in computer science from The University of Illinois at Chicago, Chicago, IL, USA, and the Ph.D. degree in information technology from the Politecnico di Milano, in 2010.

He has been an Assistant Professor with the Politecnico di Milano since 2014. His current research interests include the definition of design and analysis methodologies for embedded systems, in particular focusing on fault tolerance and reliability issues, runtime resource management in heterogeneous multi-/many-core systems, and FPGA-based systems design.



**John Lillis** (M'95) received the Ph.D. degree in computer science from the University of California at San Diego, San Diego, CA, USA, in 1996.

He is currently an Associate Professor of Computer Science with The University of Illinois at Chicago, Chicago, IL, USA, where he has been a Faculty Member since 1997. His current research interests include combinatorial optimization for Electronic Design Automation, including physical design and logic synthesis.



**Marco Domenico Santambrogio** (SM'05) received the Laurea (M.Sc. equivalent) degree in computer engineering from the Politecnico di Milano, Milan, Italy, in 2004, the M.Sc. degree in computer science from The University of Illinois at Chicago, Chicago, IL, USA, in 2005, and the Ph.D. degree in computer engineering from the Politecnico di Milano, in 2008.

He was a Post-Doctoral Fellow with the Computer Science and Artificial Intelligence Laboratory, Massachusetts Institute of Technology, Cambridge, MA, USA. He held visiting positions with the Department of Electrical Engineering and Computer Science, Northwestern University, Evanston, IL, USA, in 2006 and 2007, and the Heinz Nixdorf Institut, Paderborn, Germany, in 2006. He has been with the NECST Laboratory, Politecnico di Milano, where he founded the Dynamic Reconfigurability in Embedded System Design project in 2004 and the CHANGE (self-adaptive computing system) project in 2010. He is currently an Assistant Professor with the Politecnico di Milano. His current research interests include reconfigurable computing, self-aware and autonomic systems, hardware/software co-design, embedded systems, and high performance processors and systems.

Dr. Santambrogio is a Senior Member of the Association for Computing Machinery.